

# Lecture 5 pre-video

## Vector norms

How we measure distances  
between vectors


# Vector norms

- Named vector norms L1, L2, ... named after mathematician Henri Lebesgue (1875-1941)
- A vector norm  $p : X \mapsto \mathbb{R}$  has the following properties (where  $X$  is a vector space)
  - Triangle inequality  $p(x + y) \leq p(x) + p(y) \quad \forall x, y \in X$
  - Absolute homogeneity  $p(sx) = |s| p(x) \quad \forall s \in \mathbb{R}, x \in X$
  - Positive definiteness  $p(x) = 0 \iff x = 0$
  - Non-negativity  $p(x) \geq 0 \quad \forall x \in X$

# Vector norms

## L2: Euclidean


$$\|\mathbf{x}\|_2 \Rightarrow \sqrt{x_1^2 + x_2^2 + \dots + x_n^2}$$



# Vector norms

## L1: Absolute value / Manhattan distance

$$\| \cdot \|_1 = |x_1| + |x_2| + \dots + |x_n|$$



need abs otherwise violate  
non-negativity for vectors  
in lower/left quadrants

# Vector norms

## Application of L1 to an error function, and its derivative

For a linear fct  $f(x)$

$$\|f(x)\|_1$$



$$\frac{\partial \|f(x)\|_1}{\partial x} = \text{Sign}(f(x)) \frac{\partial f(x)}{\partial x}$$

# Vector norms

$L_\infty$


$$\|x\|_\infty = \max [x_1, x_2, \dots, x_n]$$

# Vector norms

$L_p$

$$\|x\|_p = \left( x_1^p + x_2^p + \dots + x_n^p \right)^{1/p}$$

for the  $x \in \mathbb{R}^2$  where  $x_1, x_2$  are in  
range  $[0, 1]$



# Regularization to prevent overfitting + robust regression to minimize outliers

**Jason G. Fleischer, Ph.D.**

Asst. Teaching Professor

Department of Cognitive Science, UC San Diego

[jfleischer@ucsd.edu](mailto:jfleischer@ucsd.edu)





@jasongfleischer

<https://jgfleischer.com>





# Bias Variance Tradeoff







# Intuition

## Bias-variance tradeoff in model complexity



Underfitting

(Model is too simple!)







Overfitting

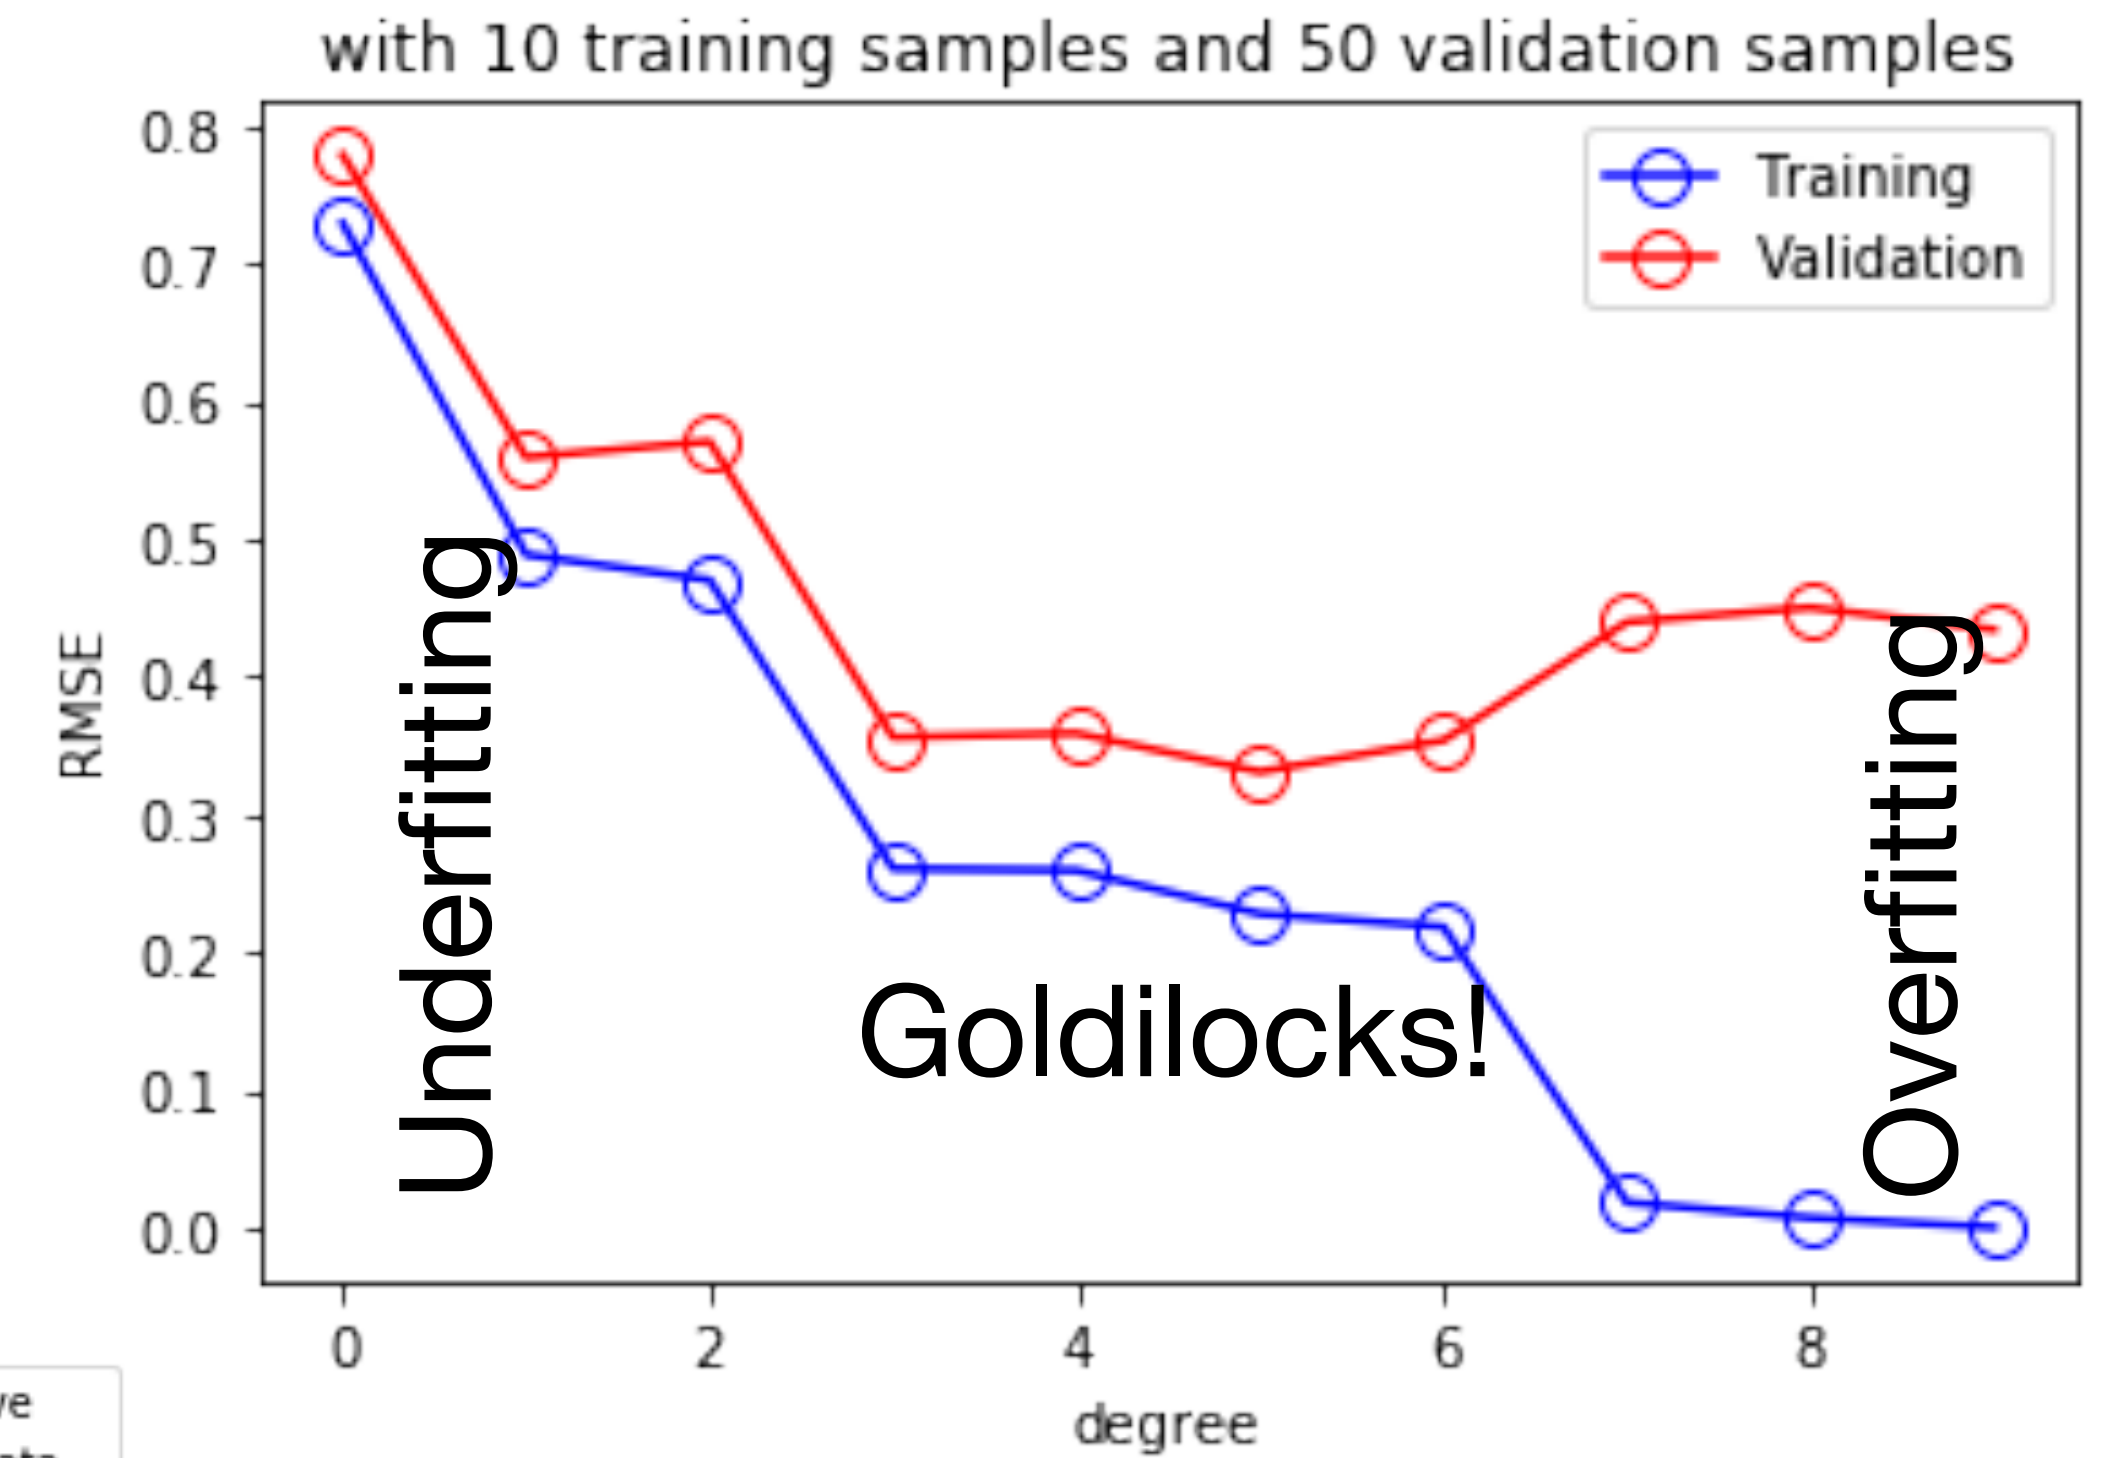
(Model is too complex!)

<b>Training set</b>	<b>-&gt; set the parameters</b>
<b>Validation set</b>	<b>-&gt; try different models, select best</b>
<b>Test set</b>	<b>-&gt; how good is your chosen model</b>

# Underfitting




Overfitting



A “validation curve”





# Implementation

## Linear Regression with Regularization



[https://colab.research.google.com/github/COGS118A/  
demo\\_notebooks/blob/main/lecture\\_06\\_regularization.ipynb](https://colab.research.google.com/github/COGS118A/demo_notebooks/blob/main/lecture_06_regularization.ipynb)

[https://github.com/COGS118A/demo\\_notebooks.git](https://github.com/COGS118A/demo_notebooks.git)

# Regularizations reduce overfitting to random noise

## Penalize solutions that are too complex

One technique that is often used to control the over-fitting phenomenon in such cases is that of **regularization**, which involves adding a penalty term to the error function below in order to discourage the coefficients from reaching large values. By preventing the sum of our weights from growing large, we are preventing complex fitting... the total amount of weight allowed will be preferentially allocated to the most important features, preventing features that have less effect on the answer from getting too much love from the algorithm.

The simplest such penalty term takes the form of a sum of squares of all of the coefficients, leading to a modified loss/error function of the form

$$L(\mathbf{w}) = (X\mathbf{w} - \mathbf{y})^T(X\mathbf{w} - \mathbf{y}) + \frac{\lambda}{2} \|\mathbf{w}\|_2^2$$

where the coefficient  $\lambda$  governs the relative importance of the regularization term compared with the sum-of-squares error term and  $\mathbf{X}$  is the (nxm) design matrix and  $\mathbf{w}$  is an m long column vector and  $\mathbf{y}$  is an n long column vector.

There is a closed-form solution below:

$$\mathbf{w}^* = (X^T X + \lambda I)^{-1} X^T \mathbf{y}$$

# Regularizations reduce overfitting to random noise

## Penalize solutions that are too complex

You can also do L1 regularization which is often called LASSO regression (least absolute shrinkage and selection operator)

$$L(\mathbf{w}) = (X\mathbf{w} - \mathbf{y})^T(X\mathbf{w} - \mathbf{y}) + \frac{\lambda}{2}\|\mathbf{w}\|_1$$


And you can combine the two together in a technique called ElasticNet

$$L(\mathbf{w}) = (X\mathbf{w} - \mathbf{y})^T(X\mathbf{w} - \mathbf{y}) + \frac{\lambda}{2}(\alpha\|\mathbf{w}\|_1 + (1 - \alpha)\|\mathbf{w}\|_2)$$

where  $\alpha \in [0, 1]$  is a parameter dictating the proportion of L1 to L2 regularization.

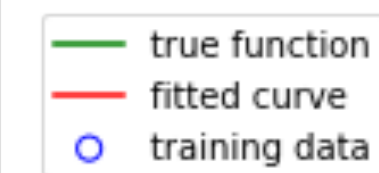








Why all these different kinds of regularization? Well L1 tends to produce a *sparse* solution... many weights that are not important are driven towards 0. You use this technique when it seems appropriate to you that less-important factors have no influence on the solution. Whereas L2 limits the total amount of weight evenly, so less-important factors can continue to have less-influence-but-still-some-influence.

One application of L1 regularization: feature selection. Let's say you have data about the expression levels of ~27,000 protein coding genes across a few thousand humans, and you want to determine if any of the genes have an effect on a disease. Since almost NONE of them will, it's good to use something like a heavy L1 penalty, which will prevent you from picking up too much on random chance associations that may exist in the data.





## 9th order polynomial regression + L2 regularization





---

# Robust estimation to reject outliers


# Estimation and optimization

---

You can have a different loss function than Residual Sum of Squares!

L2 norm: aka RSS!


$$e = \sum_{i=1}^n (y_i - f(\mathbf{x}_i; \theta))^2$$



L2

L1 norm:

$$e = \sum_{i=1}^n |y_i - f(\mathbf{x}_i; \theta)|$$



L1

# Estimation and optimization



---

$$S_{training} = \{(\mathbf{x}_i, y_i), i = 1..n\}$$

$$\mathbf{w}^* = \arg \min_{\mathbf{w}} \mathcal{L}(\mathbf{w}; \mathbf{x}; \mathbf{y})$$

$$\begin{aligned} \mathcal{L}_{OLS}(\mathbf{w}; \mathbf{x}; \mathbf{y}) &= (\mathbf{y} - f(\mathbf{x}; \mathbf{w}))^T (\mathbf{y} - f(\mathbf{x}; \mathbf{w})) \\ &= \|\mathbf{y} - f(\mathbf{x}; \mathbf{w})\|_2^2 \end{aligned}$$


$$\begin{aligned} \mathcal{L}_{robust}(\mathbf{w}; \mathbf{x}; \mathbf{y}) &= \sum_{i=1}^n |y_i - f(\mathbf{x}_i; \mathbf{w})| \\ &= \|\mathbf{y} - f(\mathbf{x}; \mathbf{w})\|_1 \end{aligned}$$



# Robust estimation


$$Y = XW$$

$$\begin{pmatrix} 1 \\ 1.9 \\ 1.05 \\ 4.1 \\ 2.1 \\ 6.0 \end{pmatrix} = \begin{pmatrix} 1, 1 \\ 1, 3 \\ 1, 2 \\ 1, 5 \\ 1, 4 \\ 1, 1.1 \end{pmatrix} \begin{pmatrix} w_0 \\ w_1 \end{pmatrix}$$



~~$$W^* = (X^T X)^{-1} X^T Y$$~~

$$W^* = \arg \min_W \sum_i (\mathbf{x}_i^T W - y_i)^2$$



$$W^* = \arg \min_W \sum_i |\mathbf{x}_i^T W - y_i|$$

$$\mathbf{w}^* = \arg \min_{\mathbf{w}} \mathcal{L}(\mathbf{w}; \mathbf{x}; \mathbf{y})$$

L1

$$S_{training} = \{(x_i, y_i), i = 1..n\}$$

$$\begin{aligned} \text{E.g. } S_{training} &= \{(x_i, y_i), i = 1..n\} \\ &= \{(1, 1), (3, 1.9), (2, 1.05), (5, 4.1), (4, 2.1)\} \end{aligned}$$

$$Y = \begin{pmatrix} 1 \\ 1.9 \\ 1.05 \\ 4.1 \\ 2.1 \end{pmatrix} \quad X = \begin{pmatrix} 1, 1 \\ 1, 3 \\ 1, 2 \\ 1, 5 \\ 1, 4 \end{pmatrix}$$

1. Loss (Cost) Function


$$\begin{aligned} \mathcal{L}_{\text{robust}}(\mathbf{y}; \mathbf{x}; \mathbf{w}) &= \sum_{i=1}^n \left| y_i - f(\mathbf{x}_i; \mathbf{w}) \right| \\ &= \|\mathbf{y} - f(\mathbf{x}; \mathbf{w})\|_1 \end{aligned}$$

2. Obtain the gradient

$$\frac{\partial \mathcal{L}(\mathbf{w})}{\partial \mathbf{w}} = \text{sign}(\mathbf{y} - f(\mathbf{x}; \mathbf{w})) \mathbf{x}$$

3. Update parameter W

$$\mathbf{w}_{t+1} = \mathbf{w}_t + \lambda_t \frac{\partial \mathcal{L}(\mathbf{w})}{\partial \mathbf{w}}$$



## Recap: Robust estimation

**Intuition:** It's very easy to overfit to either random noise or outliers. Help your ML algorithm reject either form of perturbation by applying robust estimation methods like using an **L1 loss function** (helps reduce the influence of outliers) or by adding **L1 or L2 regularization** to any loss function (helps reduce overfitting to noise by penalizing complex and large parameter values)

### Math:

L1 loss function  
(robust regression)


$$L(\mathbf{w}) = \frac{1}{2} \|X\mathbf{w} - \mathbf{y}\|_1$$

L2 loss function + combo L1/L2 regularization  
(ElasticNet)

$$L(\mathbf{w}) = \frac{1}{2} (X\mathbf{w} - \mathbf{y})^T (X\mathbf{w} - \mathbf{y}) + \frac{\lambda}{2} (\alpha \|\mathbf{w}\|_1 + (1 - \alpha) \|\mathbf{w}\|_2)$$

# Overfitting and underfitting

- Normally  $e_{\text{train}} \leq e_{\text{test}}$
- $e_{\text{test}} = e_{\text{train}} + e_{\text{generalization}}$
- Overfit: when  $e_{\text{train}} \ll e_{\text{test}}$ 
  - Generalization error gets big!
- Underfit:
  - $e_{\text{train}} \sim \text{chance level}$ ,
  - $e_{\text{test}}$  going down with additional model complexity






# Testing error

$$S_{training} = \{(\mathbf{x}_i, y_i), i = 1..n\} \quad S_{testing} = \{(\mathbf{x}_i, y_i), i = 1..q\}$$

$$e_{testing} = e_{training} + generalization(f)$$



$$e_{testing} = \frac{1}{q} \sum_{i=1}^q \mathbf{1}(y_i \neq f(\mathbf{x}_i))$$

$$e_{testing} = 0.5?$$

$$1. e_{testing} = 0.5 + 0.0$$

$$2. e_{testing} = 0.0 + 0.5$$

Question: What are the two extreme cases leading to the same testing error=0.5 (Assume we have the same number of positives and negatives.)



# Testing error

---

$$S_{training} = \{(\mathbf{x}_i, y_i), i = 1..n\} \quad S_{testing} = \{(\mathbf{x}_i, y_i), i = 1..q\}$$

$$\epsilon_{testing} = \epsilon_{training} + \epsilon_{generalization}$$

## Case 1:

$$1. \epsilon_{testing} = 0.5 + 0.0$$

- We make a **completely random** guess even after training (didn't learn anything and attained an error of 0.5).
- A random guess is however **highly generalizable**, which doesn't incur any generalization error.

## Case 2:

$$2. \epsilon_{testing} = 0.0 + 0.5$$

- We make perfect classifications on the training data (memorizing the labels for every training sample).
- Merely memorizing the training samples with their corresponding classification labels is however **highly non-generalizable**, incurring the largest generalization error (no identical training sample will appear in testing).

Both are extreme cases that lead to trivial ML models.

# Testing error

---

$$\epsilon_{\text{testing}} = \epsilon_{\text{training}} + \epsilon_{\text{generalization}}$$

Case 1:

1.  $e_{\text{testing}} = 0.5 + 0.0$

Case 2:

2.  $e_{\text{testing}} = 0.0 + 0.5$

Both are extreme cases that lead to trivial models that we should avoid.

**Ideally:**  $e_{\text{testing}} = 0.0 + 0.0$

A classification model that is perfect after training and makes no error in testing.

---

**In practice:**  $e_{\text{testing}} = 0.05 + 0.1$

A classification model that does well after training and generalizes reasonably well in testing.