

Model selection

Jason G. Fleischer, Ph.D.

Asst. Teaching Professor

Department of Cognitive Science, UC San Diego

jfleischer@ucsd.edu



@jasongfleischer

<https://jgfleischer.com>


Slides in this presentation are from material kindly provided by
Sebastian Rashka



Sample p : 0.67




Sample p : 0.56




Sample p : 0.55




Sample p : 0.64




Sample p : 0.58




Sample p : 0.57



Sample p : 0.59



Sample p : 0.63



Sample p : 0.67



Sample p : 0.58

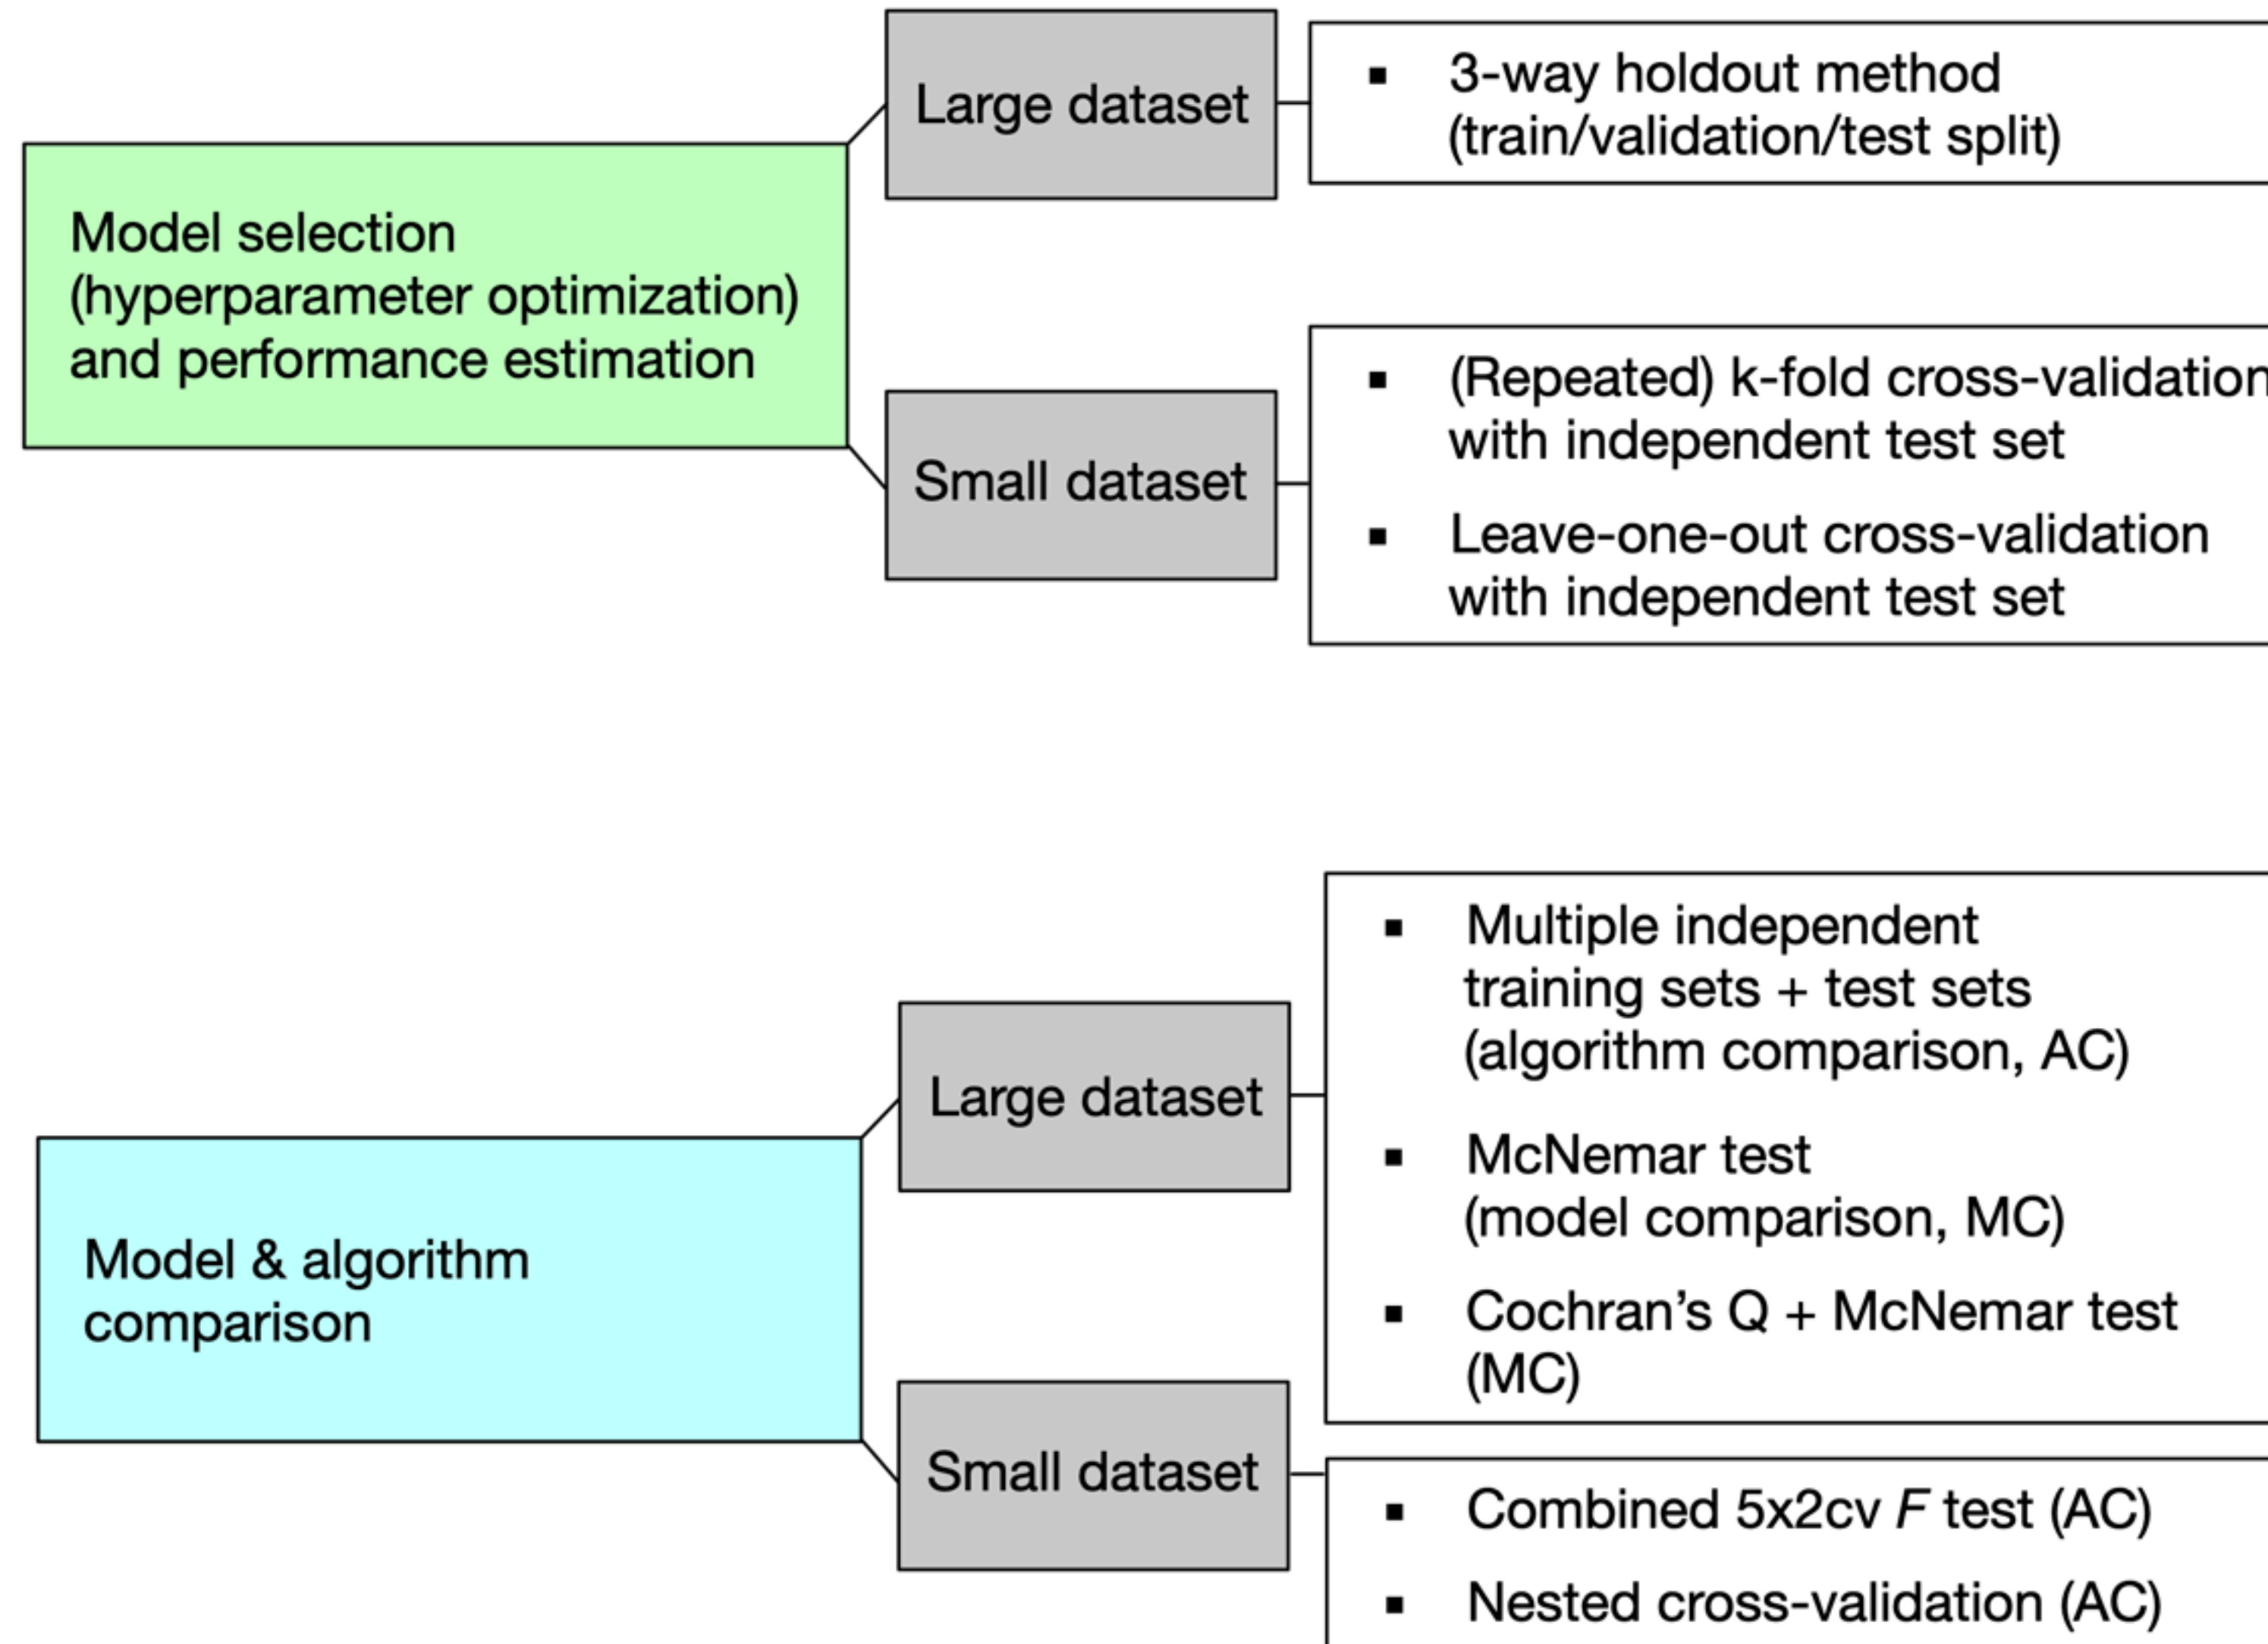
One of the
hold out folds

Or one trial of
model selection
via cross-val

Estimation of performance

Many methods, two use cases, one reason

- **THE ONE REASON:** every measure is a random draw from a distribution of performances... What if the data was a bit different? What if the random seed is different? Etc.
- **TWO USE CASES:**
 - To estimate how well the system will generalize (test)
 - To perform model selection or algorithm selection (validation)
- **MANY METHODS:**
 - See Model Evaluation, Model Selection, and Algorithm Selection in Machine Learning
Sebastian Rashkha
<https://arxiv.org/pdf/1811.12808.pdf> for a good intro (but there are more out there!)



Loss function

Parameters
e.g., weight vector

Algorithm
e.g., Logistic Regression

Model

Literal algorithm
e.g. prediction
function, training
method, etc

Hyper-parameters
e.g., regularization setup, solver

Loss function

Model 1

Model 2

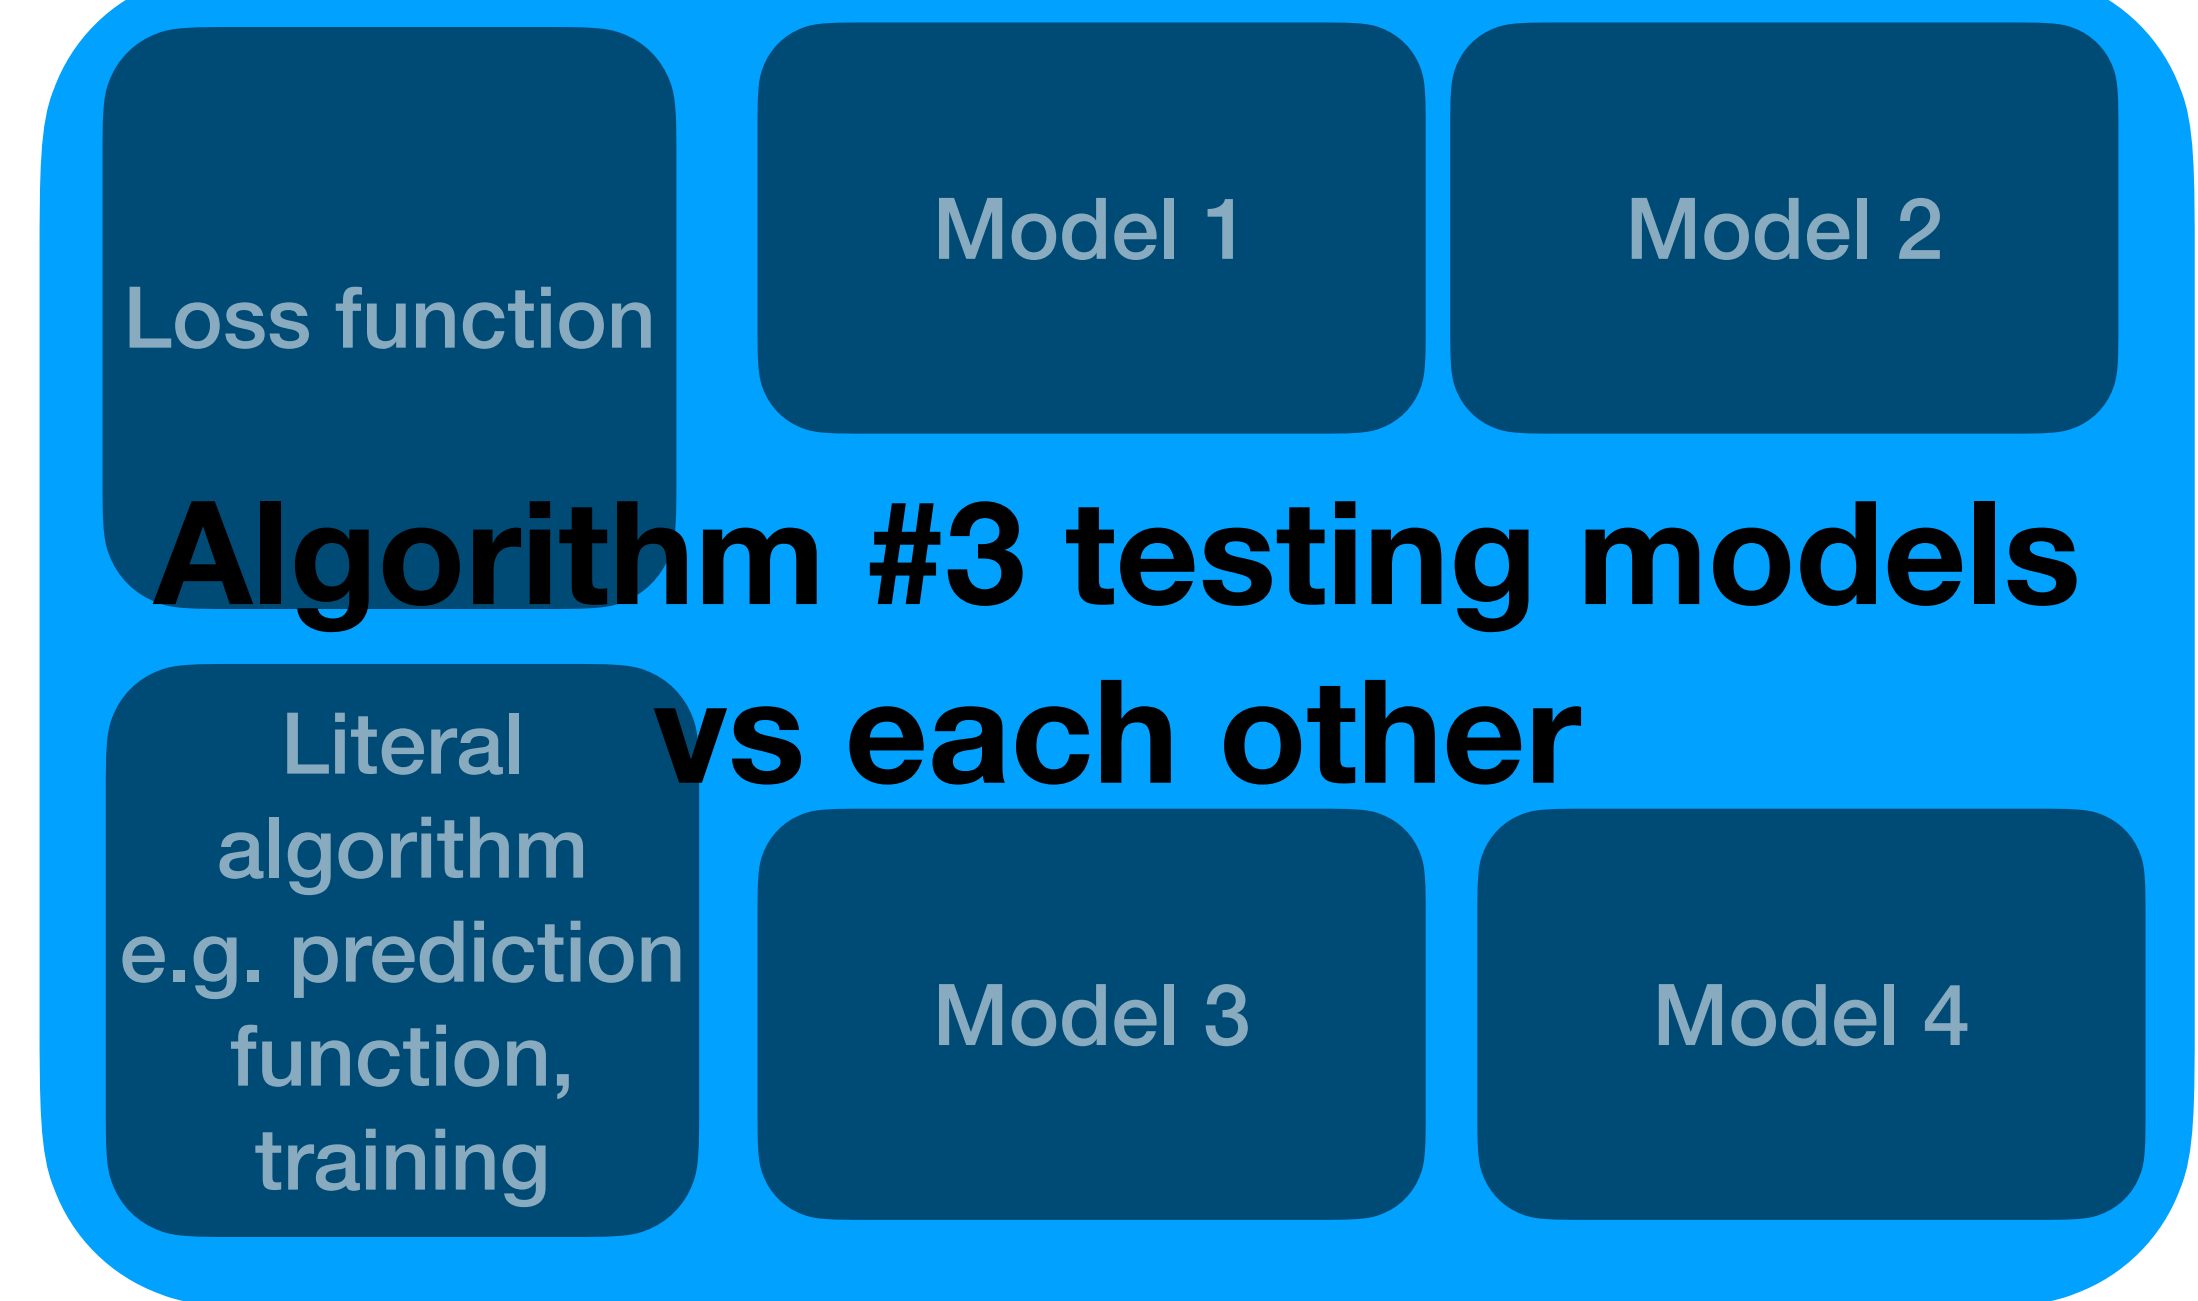
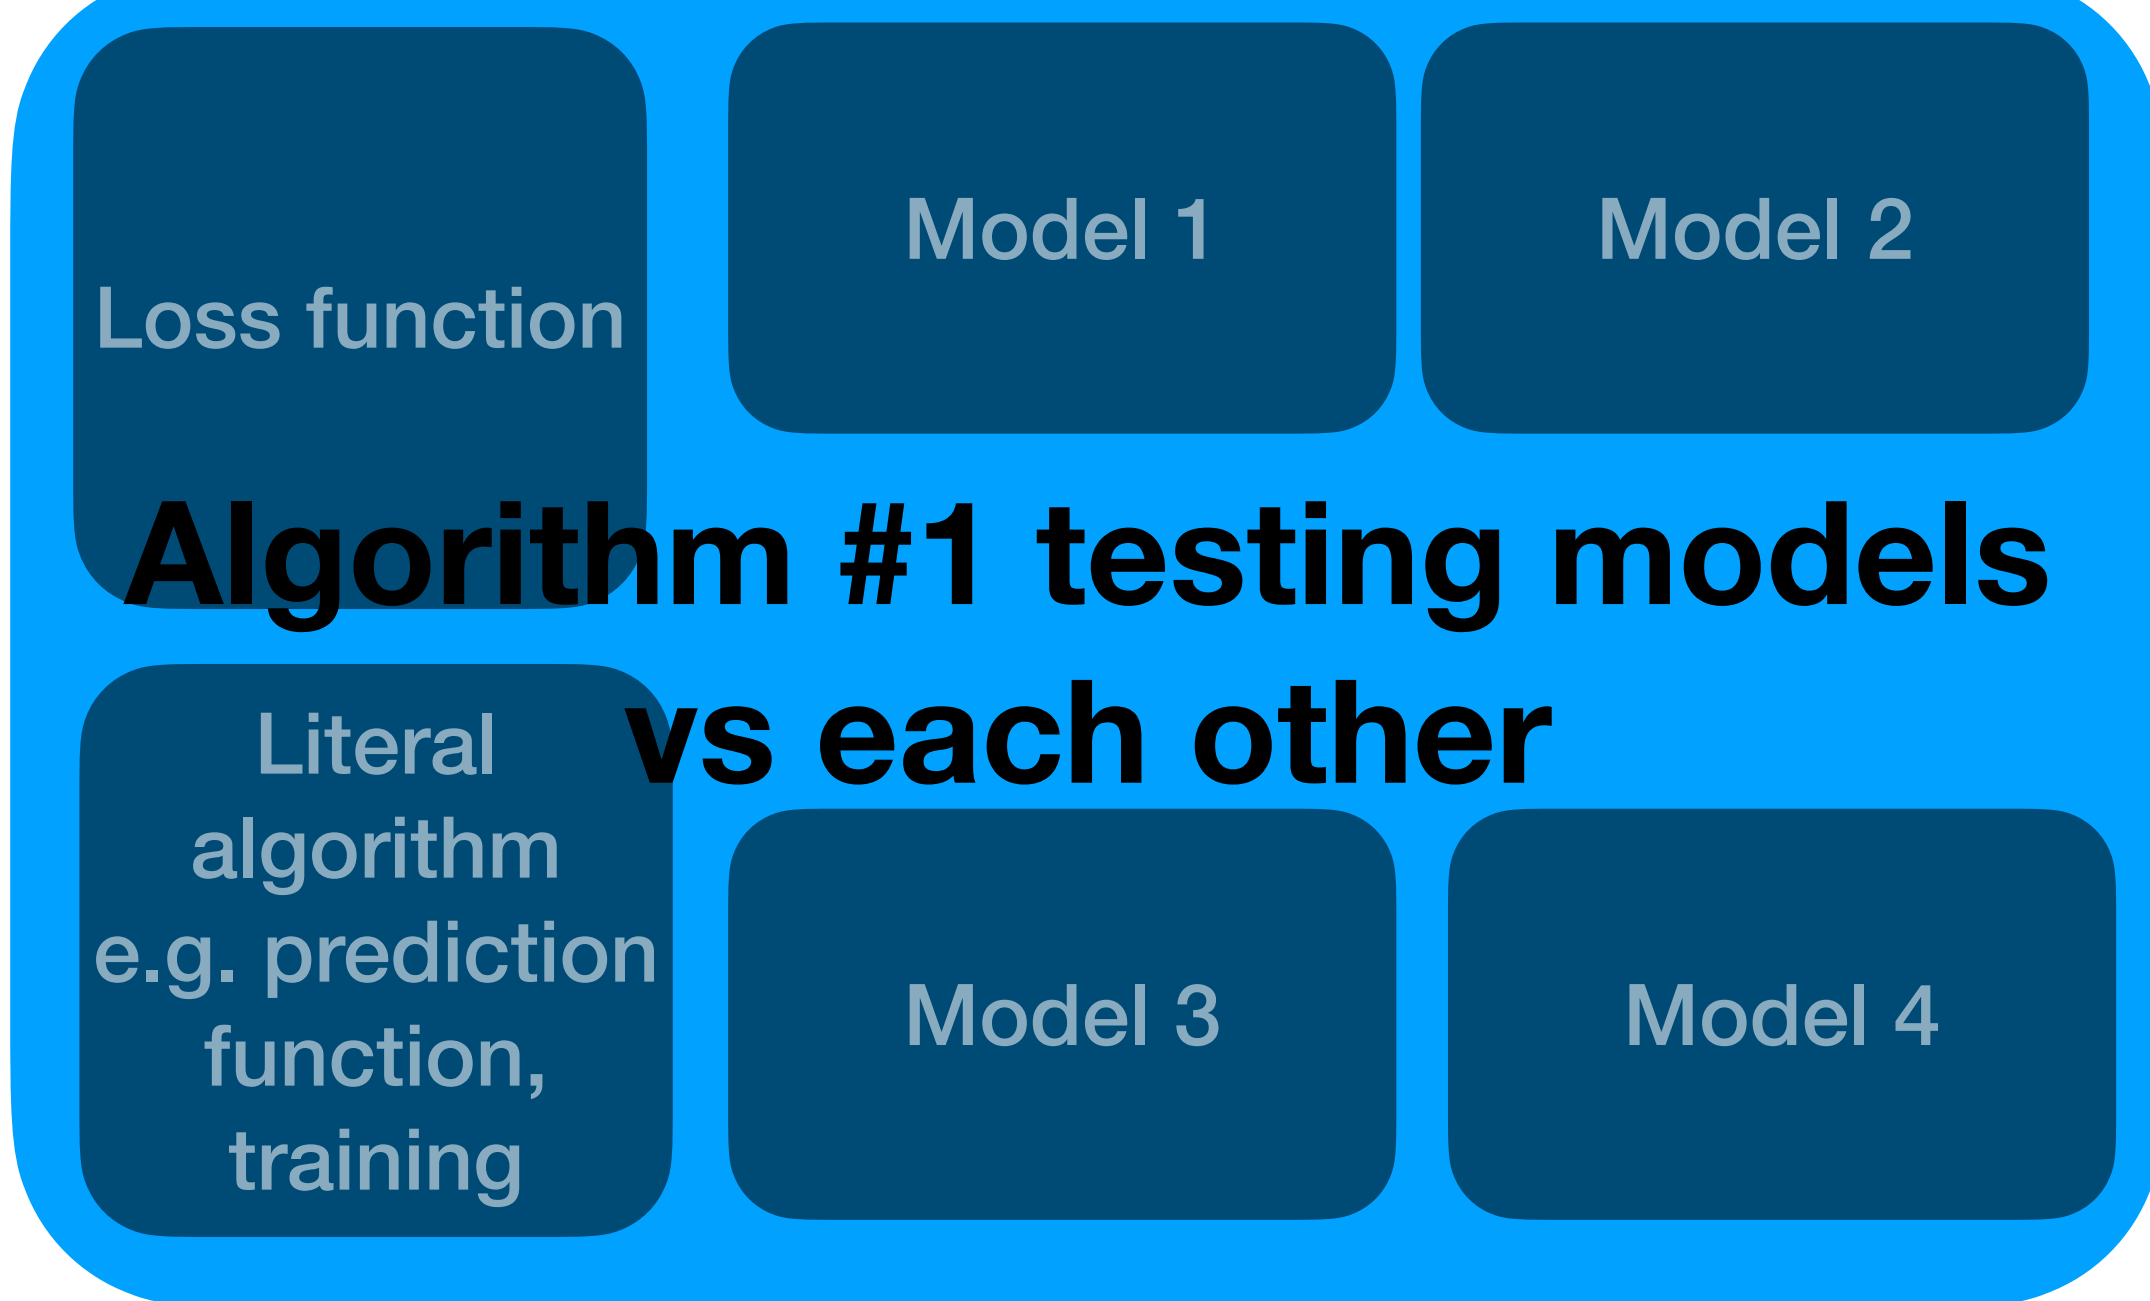
**Single algorithm testing models
vs each other**

Literal algorithm
e.g. prediction
function, training
method, etc

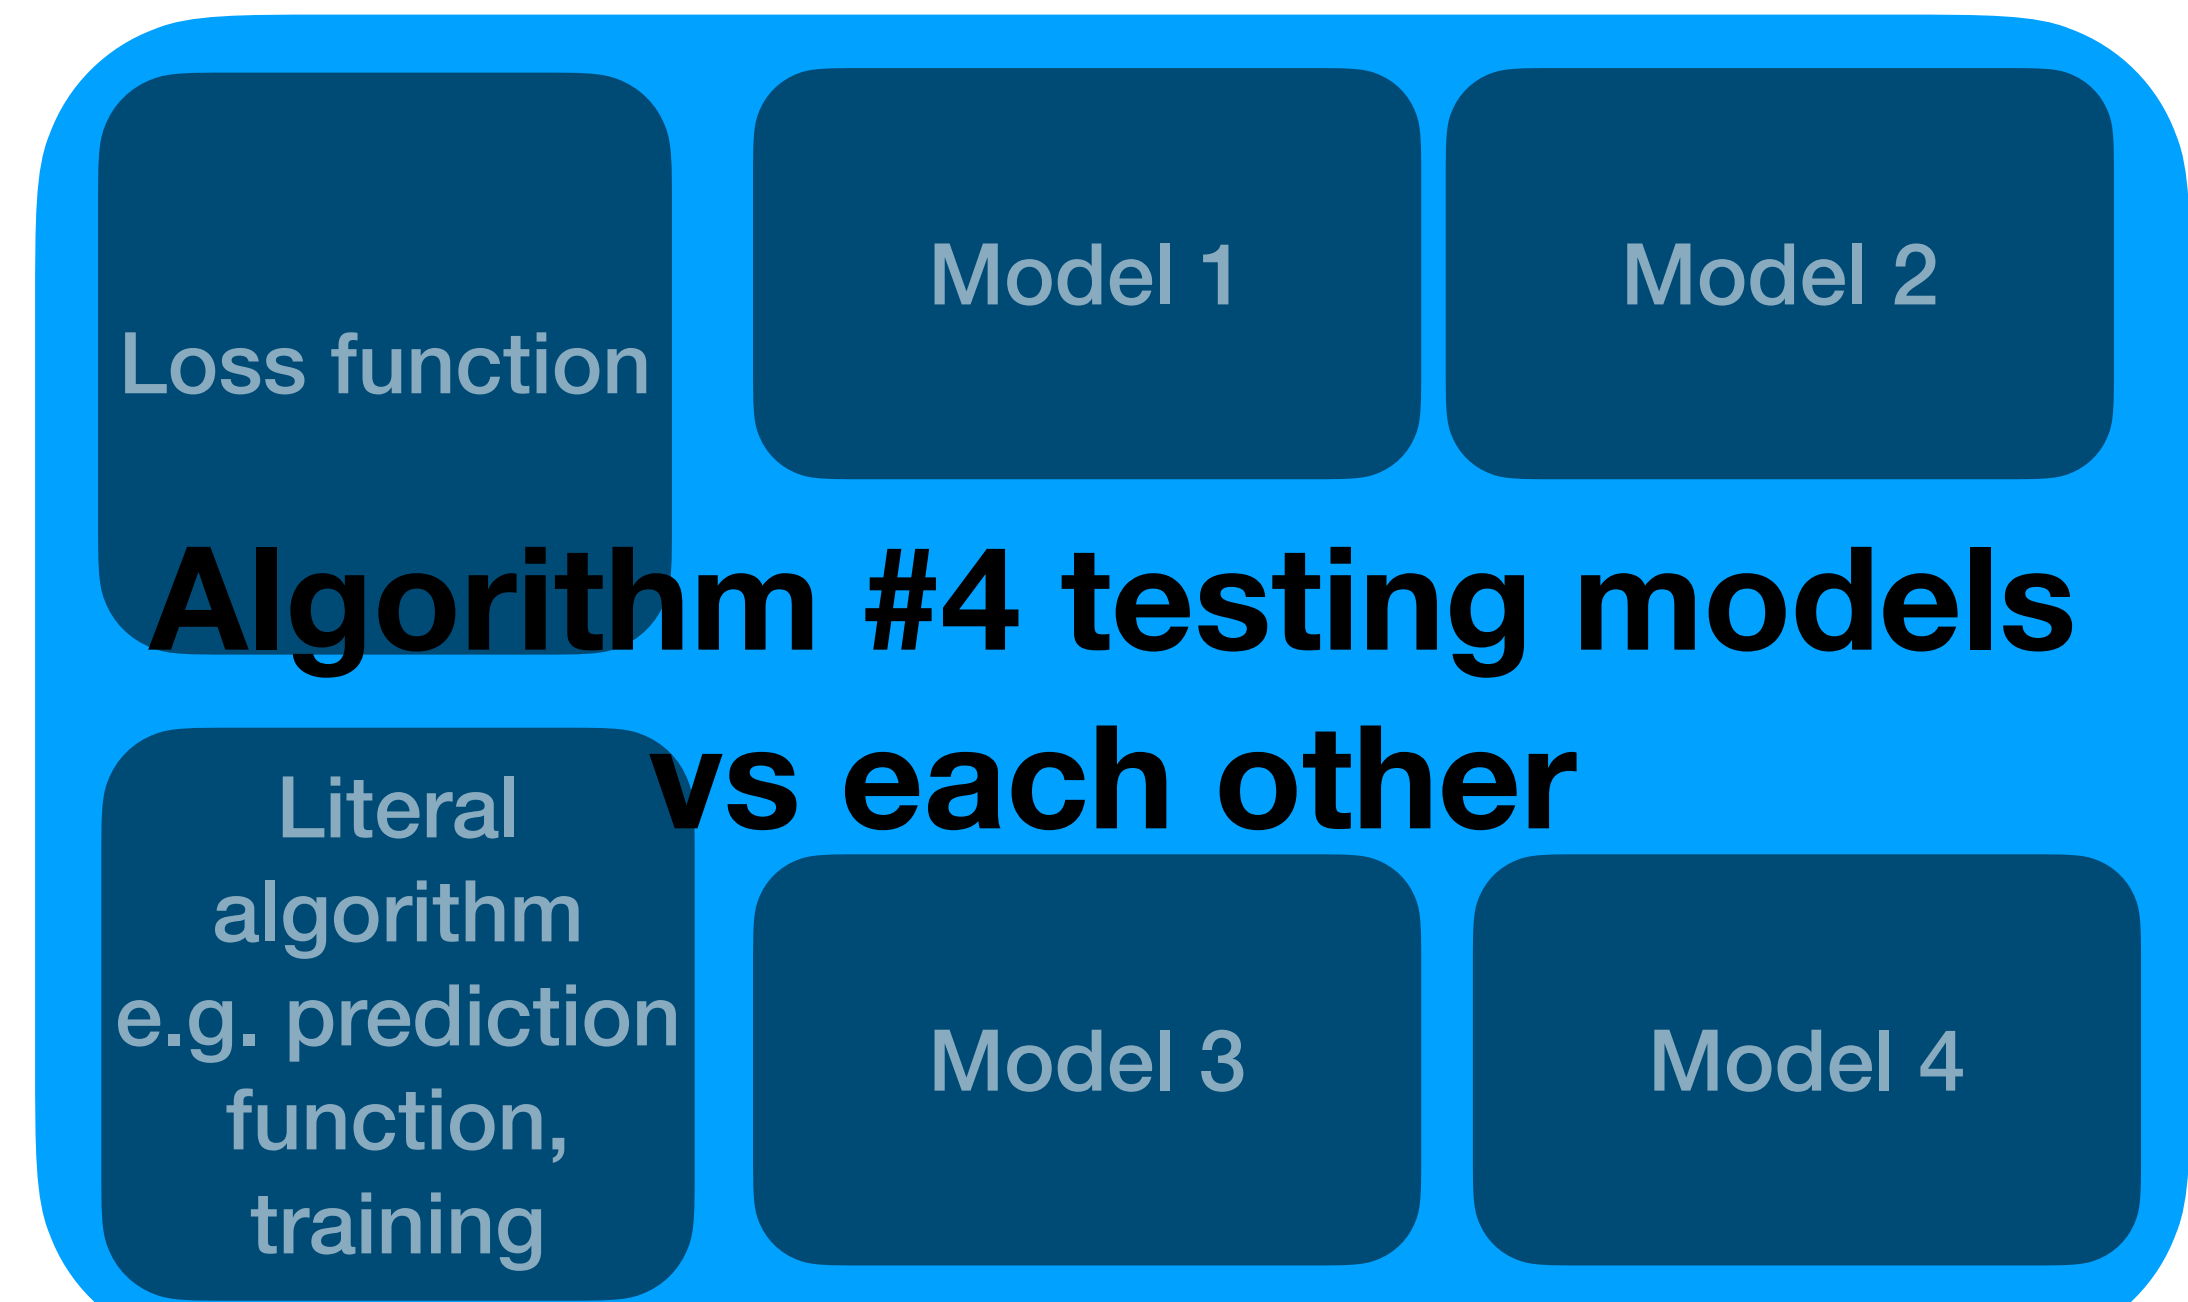
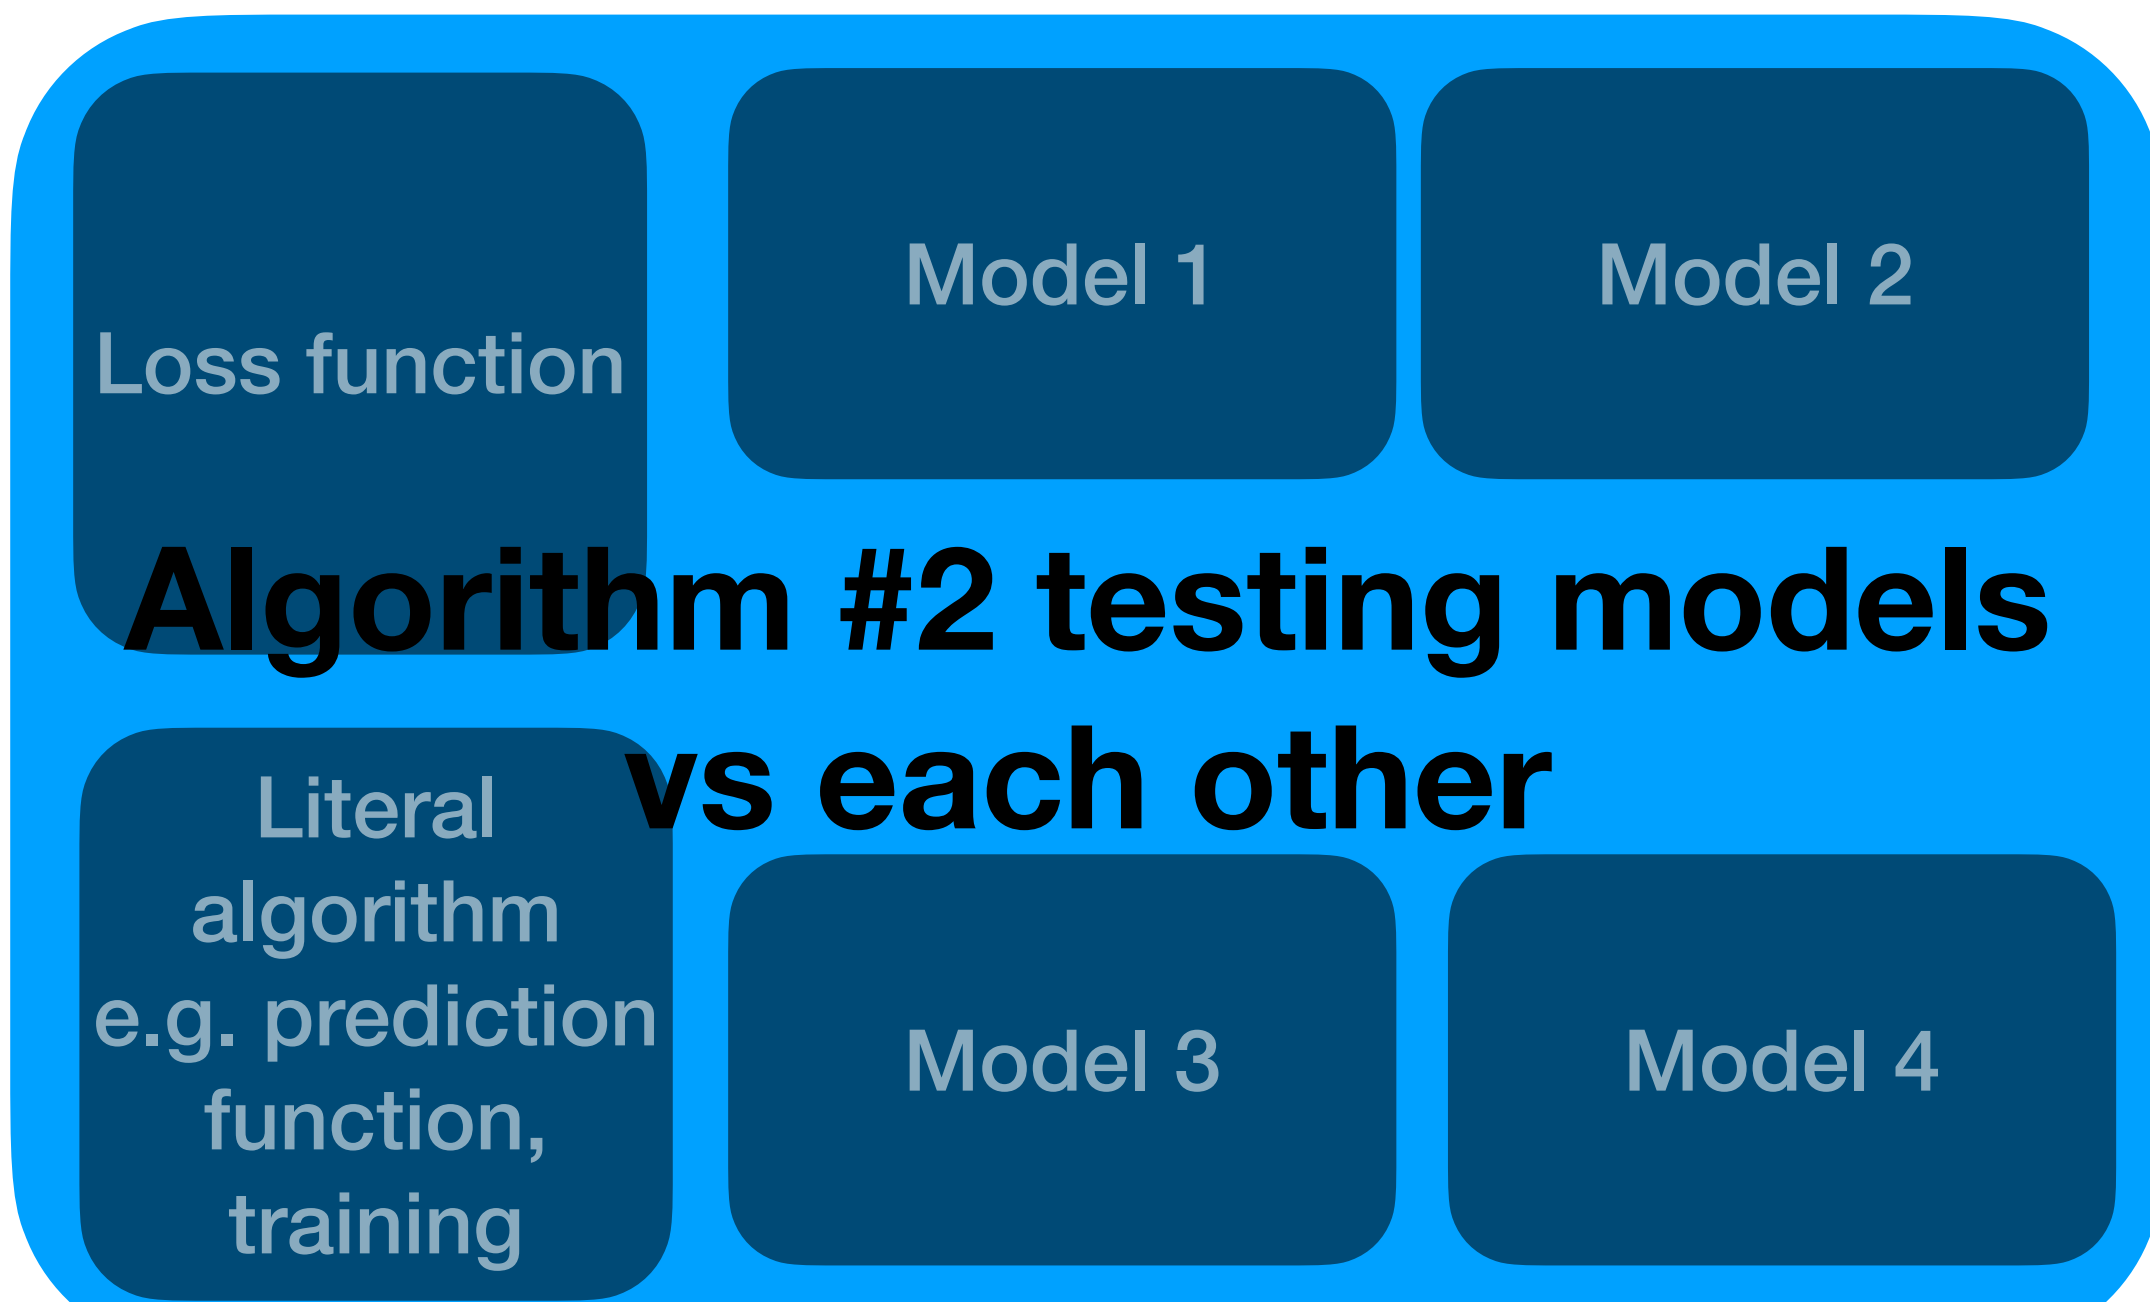
Model 3

Model 4

“Model selection”



“Algorithm selection”



Method #1 - Train/Validate/Test sets

For either Model or Algorithm selection using HUGE datasets


- Split data into train, validate, test
- [OPTIONAL] Outer loop... do this T times:
 - do this M times, once for each model in the hyper-parameter search space or each algorithm-model combination:
 - Train it on the same training set
 - Predict on the same validation set
- Pick the best model or algorithm based on its performance on [OPTIONAL the mean across trials] of the validation set
- Train the best version on the whole of training set + validation set
- Test it on the test set to estimate its ability to generalize

METHOD 1 - with enough data to have a good test set

Let's say you had around 8k samples in a dataset

For each trial:

- training set ~ sample 5k (with or w/o) replacement from entire dataset
- Grid search of hyper parameters using k-fold cross validation on the training set
- Select best model from grid, train on entire training set
- Evaluate best model on the test set (everything not sampled for training)



Method #2 - Cross validation

For either model or algorithm selection using medium sized datasets

- Split data into cross-validation and test sets
- [OPTIONAL] Outer loop... do this T times:
 - do this M times, once for each model in the hyper-parameter search space or each algorithm-model combination:
 - Use k-fold cross validation to estimate validation error
- Pick the best model or algorithm based on its performance on [OPTIONAL the mean across trials] of the validation sets
- Train the best version on the whole of cross validation set
- Test it on the test set to estimate its ability to generalize

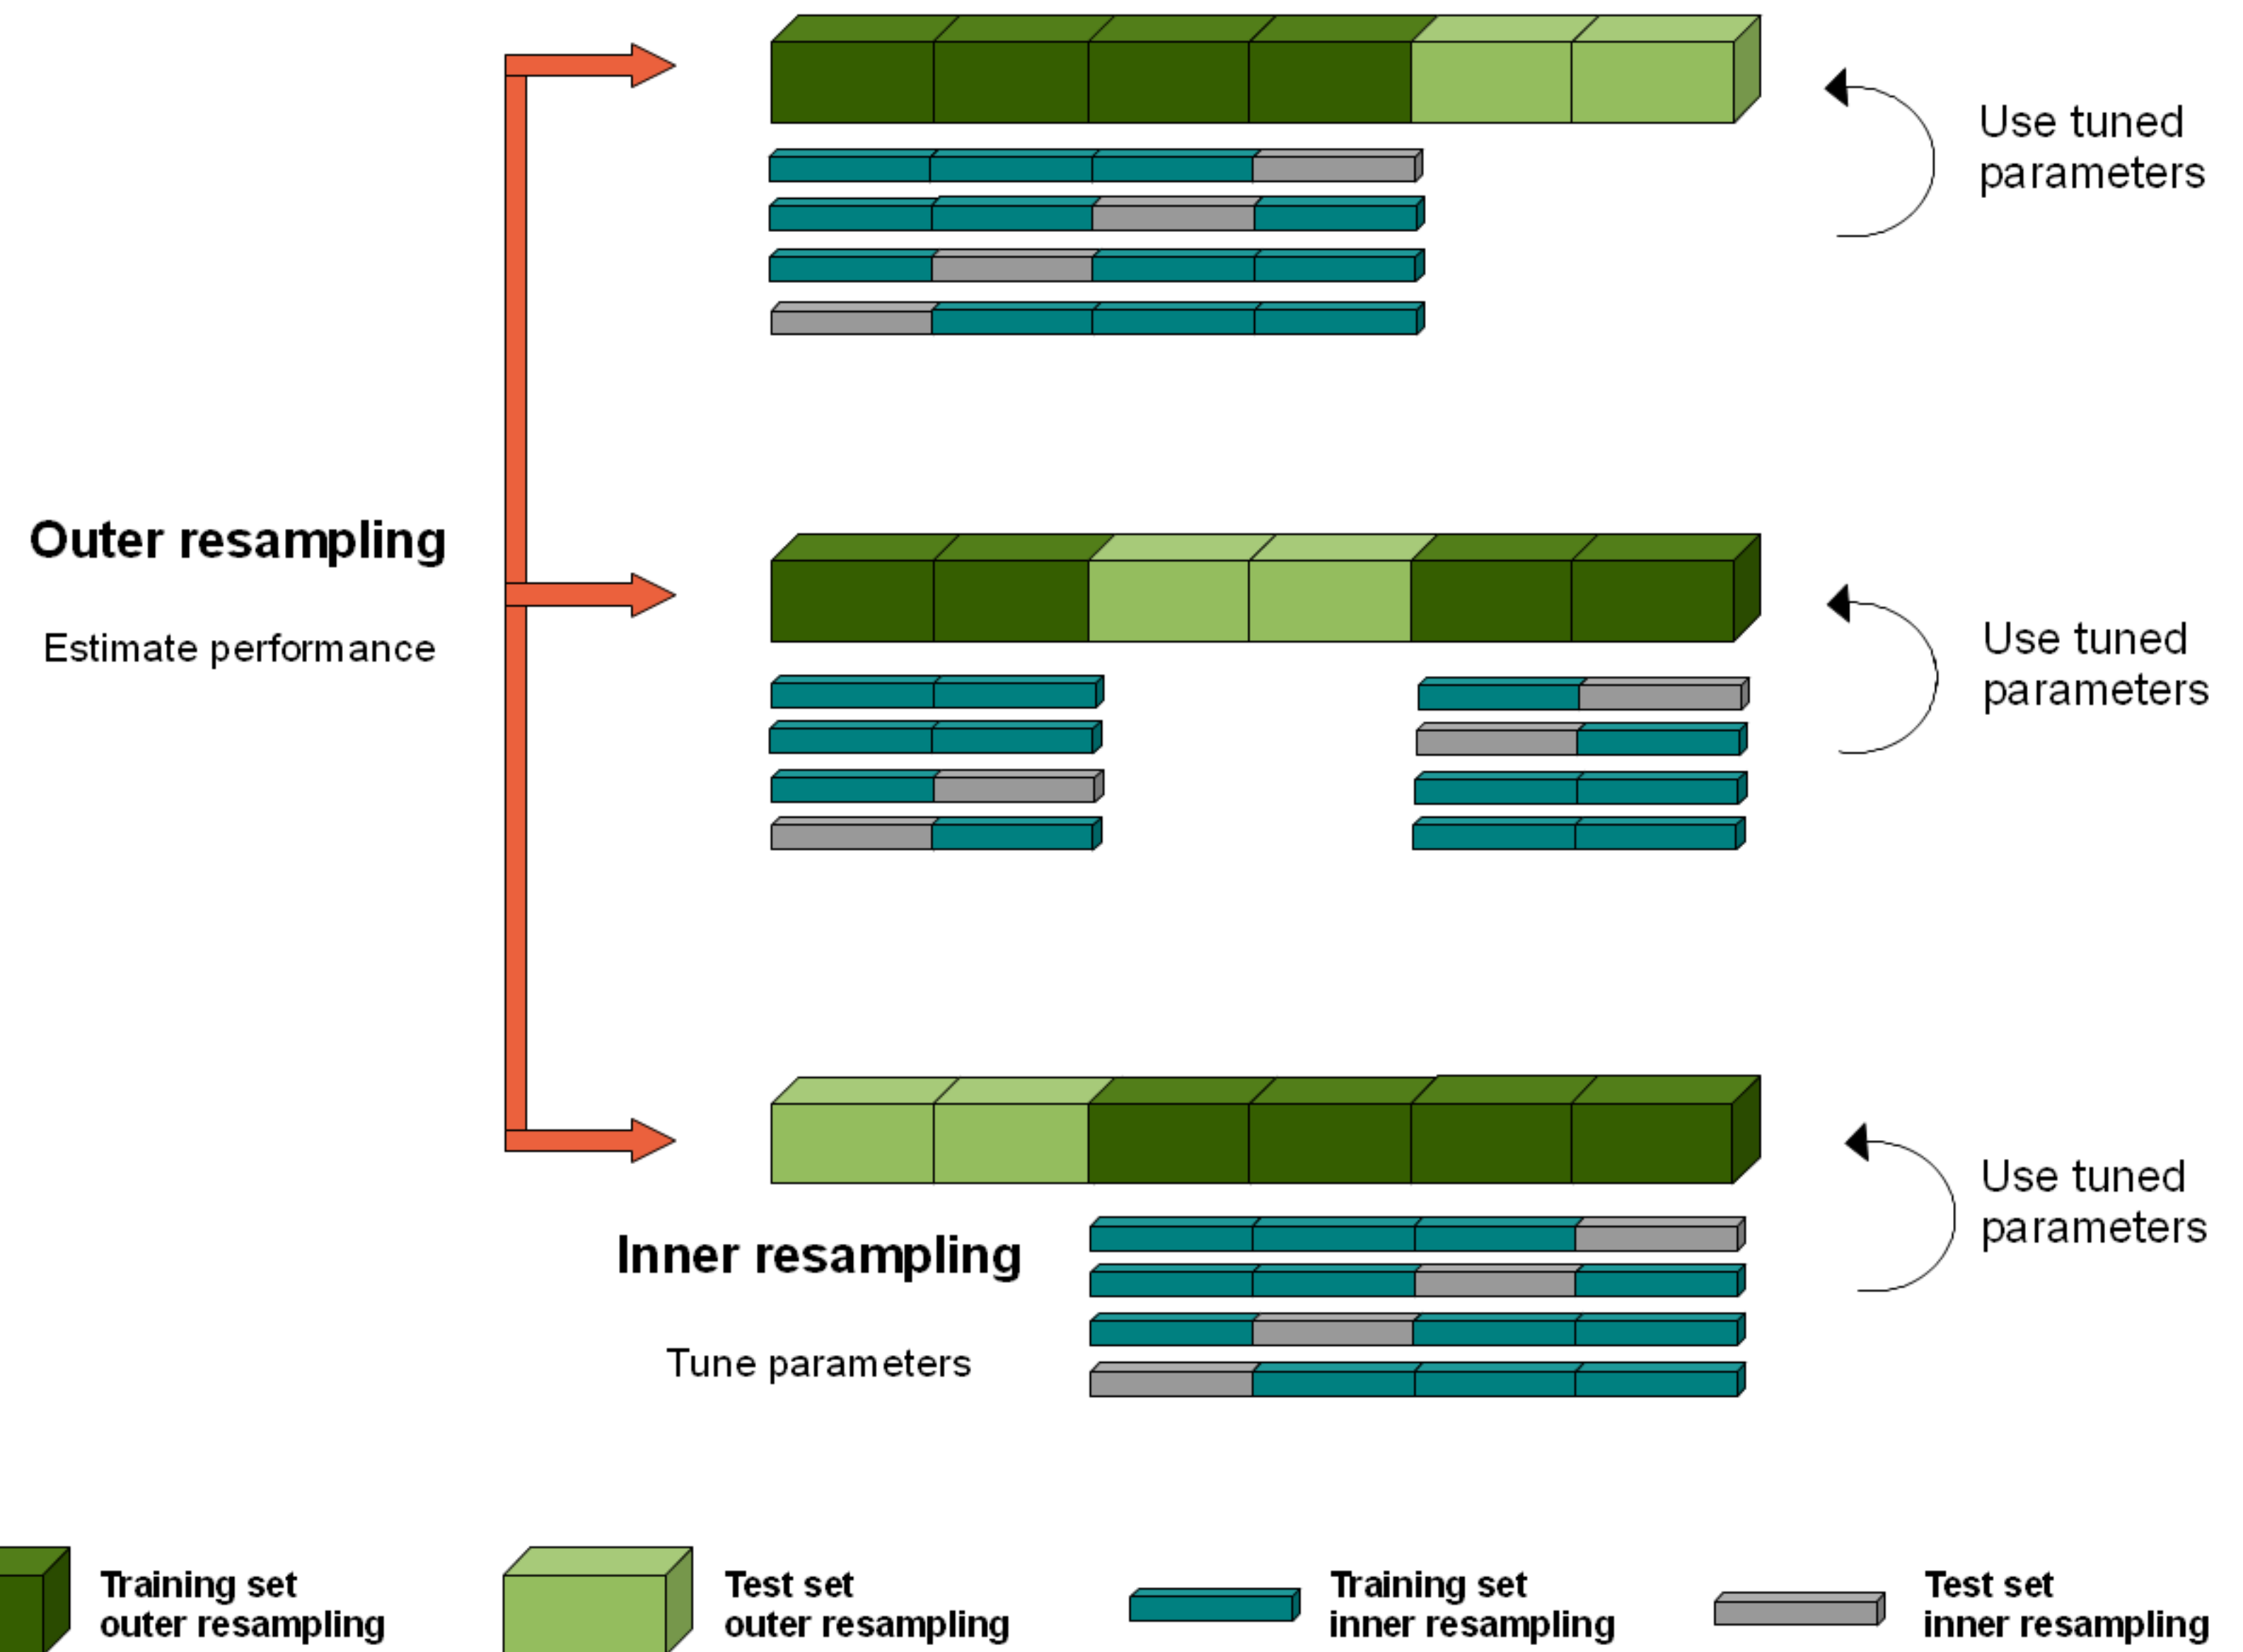
METHOD 2 - make the most of a small amount of data

Nested Cross-validation

For Algorithm Comparison if done the time efficient way...
(only best hyperparams tested on the outer cross-val)

...can be used for Model comparison if done the inefficient way (all hyper params tested on the outer cross-val)

This for when you've got only ~2000 samples, which is barely enough to fit the data well let alone test



Method #2a - Nested CV for al

For doing algorithm selection on medium sized datasets

- Do not split off a test set!
- [OPTIONAL] Outer loop... do this T times:
 - Do this M times, once for each algorithm
 - Use nested k-fold cross validation...
 - Inner loop estimates validation error for all the hyperparams tested for a given model
 - Outer loop estimates validation error for a given algorithm
- Pick the best algorithm based on its performance on [OPTIONAL the mean across trials] of the outer cross validation folds

Model selection with built-in test set error using nested CV

Method #2a - Small sized datasets

- Do not split off a test set!
- [OPTIONAL] Outer loop... do this T times:
 - Do this M times, once for each algorithm
 - Use nested k-fold cross validation...
 - Inner loop estimates validation error for all the hyperparams tested for a given model
 - Outer loop estimates validation error for a given algorithm
- Pick the best algorithm based on its performance on [OPTIONAL the mean across trials] of the outer cross validation folds


METHOD 2 - make the most of a small amount of data

Nested Cross-validation

For Algorithm Comparison if done the time efficient way...
(only best hyperparams tested on the outer cross-val)

...can be used for Model comparison if done the inefficient way (all hyper params tested on the outer cross-val)


This for when you've got only ~2000 samples, which is barely enough to fit the data well let alone test



**But how do you organize your
search of the hyper parameter
space?**

Grid Search

- Exhaustive search
- Thorough but expensive
- Specify grid for parameter search
- Can be run in parallel
- Can suffer from poor coverage
- Often run with multiple resolutions



Bergstra, J., & Bengio, Y. (2012). Random search for hyperparameter optimization. *The Journal of Machine Learning Research*, 13(1), 281-305.

Randomized Search

- Search based on a time budget
- Preferred if there are many hyperparameters (e.g. > 3 distinct ones)
- specify distribution for parameter search
- can be run in parallel




Figure 1: Grid and random search of nine trials for optimizing a function $f(x,y) = g(x) + h(y) \approx g(x)$ with low effective dimensionality. Above each square $g(x)$ is shown in green, and left of each square $h(y)$ is shown in yellow. With grid search, nine trials only test $g(x)$ in three distinct places. With random search, all nine trials explore distinct values of g . This failure of grid search is the rule rather than the exception in high dimensional hyper-parameter optimization.

Bergstra, J., & Bengio, Y. (2012). Random search for hyper-parameter optimization. *The Journal of Machine Learning Research*, 13(1), 281-305.


Statistical testing

https://sebastianraschka.com/pdf/lecture-notes/stat479fs18/11_eval-algo_notes.pdf

Statistical testing on model performance

- Testing is almost always paired (over folds of cross validation)
- Distinguish between tests appropriate for algorithm comparison vs model selection (hyperparameter settings)
- Distinguish between test that are computationally efficient vs those that are not
- Distinguish between pair-wise and group-wise tests


Image Classification on ImageNet



**Jason gets grumpy about blindly
following methods you don't
understand fully**

The p-value





- In range 0,1
- Smaller is support for alternative hypothesis
- Larger is inconclusive
- Ignores effect size!!!@!!! Is the difference practically important?
- Assumes conditions on data
- $$P(H|D) = \frac{P(D|H)P(H)}{P(D)}$$



Parsimony Principle


Choose the simplest w/in 1 std error of optimal

Which parameter would you select?



Maybe you don't need a statistical test

3x5 repeated cross validation results



No free lunch theorem

Why even bother??

